

# Final Team Project - Team 4

Uyen Pham, Jacqueline Urenda, and Renetta Nelson

December 5, 2022

## Loading Libraries

```
library(tidyverse)
library(fpp2)
library(readr)
library(forecast)
library(ggplot2)
library(gridExtra)
library(reshape2)
library(dplyr)
library(lubridate)
library(RColorBrewer)
library(corrplot)
library(Hmisc)
library(ggpubr)
set.seed(506)
```

## Loading the Dataset

```
#Load the data
power <- read_csv("Tetuan City power consumption.csv",
  col_types = cols(DateTime = col_datetime(format = "%m/%d/%Y %H:%M")))
```

The columns in the data set are renamed. This was done for easy reference to variables.

```
colnames(power) <- c('DateTime', 'Temperature', 'Humidity', 'Wind_Speed',
  'Gen_Diffuse_Flows', 'Diffuse_Flows', 'Zone1', 'Zone2',
  'Zone3')
```

```
head(power)
```

```
## # A tibble: 6 x 9
##   DateTime      Temperature Humidity Wind_~1 Gen_D~2 Diffu~3 Zone1 Zone2
##   <dtm>          <dbl>     <dbl>   <dbl>   <dbl>   <dbl>   <dbl> <dbl>
## 1 2017-01-01 00:00:00      6.56      73.8    0.083    0.051    0.119 34056. 16129.
## 2 2017-01-01 00:10:00      6.41      74.5    0.083    0.07     0.085 29815. 19375.
## 3 2017-01-01 00:20:00      6.31      74.5    0.08     0.062    0.1    29128. 19007.
## 4 2017-01-01 00:30:00      6.12      75     0.083    0.091    0.096 28229. 18361.
## 5 2017-01-01 00:40:00      5.92      75.7    0.081    0.048    0.085 27336. 17872.
## 6 2017-01-01 00:50:00      5.85      76.9    0.081    0.059    0.108 26625. 17416.
## # ... with 1 more variable: Zone3 <dbl>, and abbreviated variable names
## #   1: Wind_Speed, 2: Gen_Diffuse_Flows, 3: Diffuse_Flows
```

## Exploratory Data Analysis

```
print("Missing Values: ")

## [1] "Missing Values: "

sum(is.na(power)) #no missing values

## [1] 0
```

## Statistical Data Analysis

The data set consists of 364 days total, taking data from January 1, 2017 to December 30, 2017. The time window is every ten minutes. The temperature is measured in Celsius with a mean and median around 18 degrees. The humidity column displays the percentage of the humidity. The average humidity is 68.26%. The wind speed is measured in km/h, and the power consumption is measured in KiloWatts.

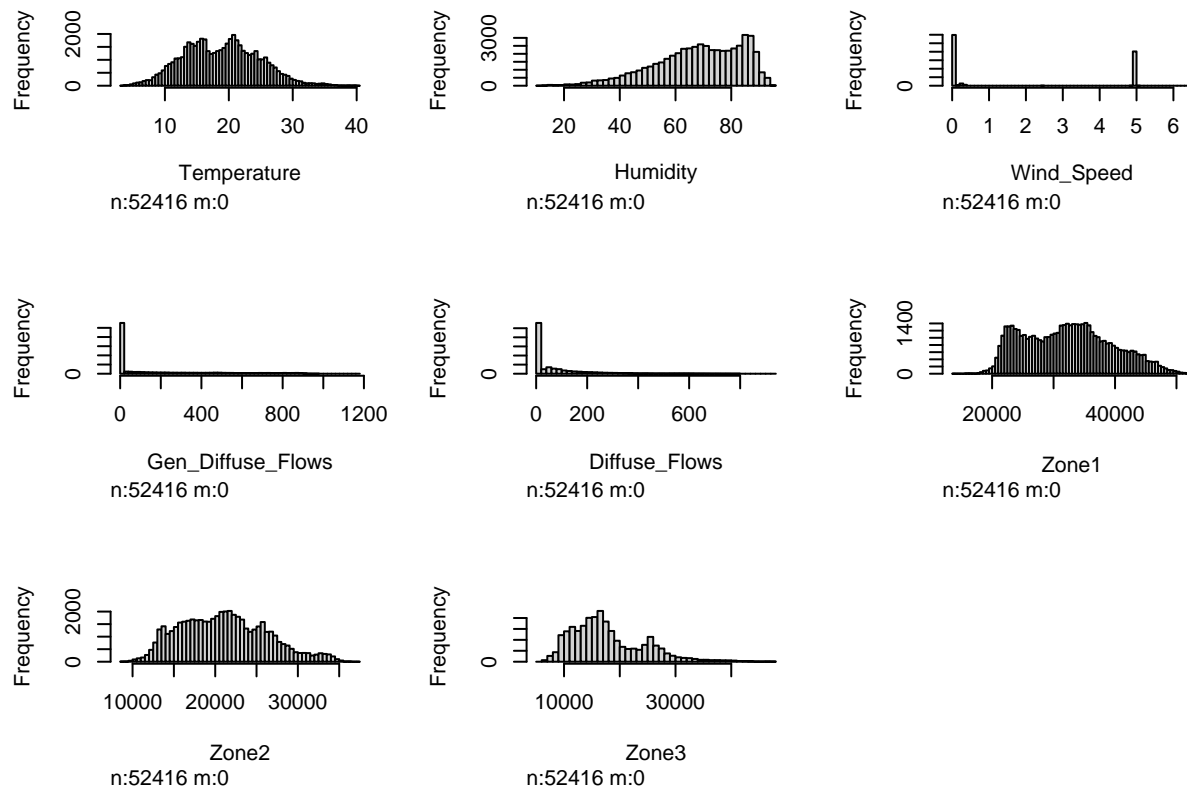
```
summary(power)
```

##	DateTime	Temperature	Humidity	Wind_Speed
##	Min. :2017-01-01 00:00:00	Min. : 3.247	Min. :11.34	Min. :0.050
##	1st Qu.:2017-04-01 23:57:30	1st Qu.:14.410	1st Qu.:58.31	1st Qu.:0.078
##	Median :2017-07-01 23:55:00	Median :18.780	Median :69.86	Median :0.086
##	Mean :2017-07-01 23:55:00	Mean :18.810	Mean :68.26	Mean :1.959
##	3rd Qu.:2017-09-30 23:52:30	3rd Qu.:22.890	3rd Qu.:81.40	3rd Qu.:4.915
##	Max. :2017-12-30 23:50:00	Max. :40.010	Max. :94.80	Max. :6.483
##	Gen_Diffuse_Flows	Diffuse_Flows	Zone1	Zone2
##	Min. : 0.004	Min. : 0.011	Min. :13896	Min. : 8560
##	1st Qu.: 0.062	1st Qu.: 0.122	1st Qu.:26311	1st Qu.:16981
##	Median : 5.035	Median : 4.456	Median :32266	Median :20823
##	Mean :182.697	Mean : 75.028	Mean :32345	Mean :21043
##	3rd Qu.:319.600	3rd Qu.:101.000	3rd Qu.:37309	3rd Qu.:24714
##	Max. :1163.000	Max. :936.000	Max. :52204	Max. :37409
##	Zone3			
##	Min. : 5935			
##	1st Qu.:13129			
##	Median :16415			
##	Mean :17835			
##	3rd Qu.:21624			
##	Max. :47598			

## Histograms

Temperature appears to have a normal distribution. Humidity is left skewed- as mentioned in summary statistics, humidity is pretty high. Wind Speed is interesting as well, looks like either little to low wind speeds or ~4.5 km/h wind speed. General and Diffuse flows are right skewed. Flow tends to be on a lower level. Zone 1 and Zone 2 have a normal distribution. Zone 2 looks normal as well but looks a bit right skewed.

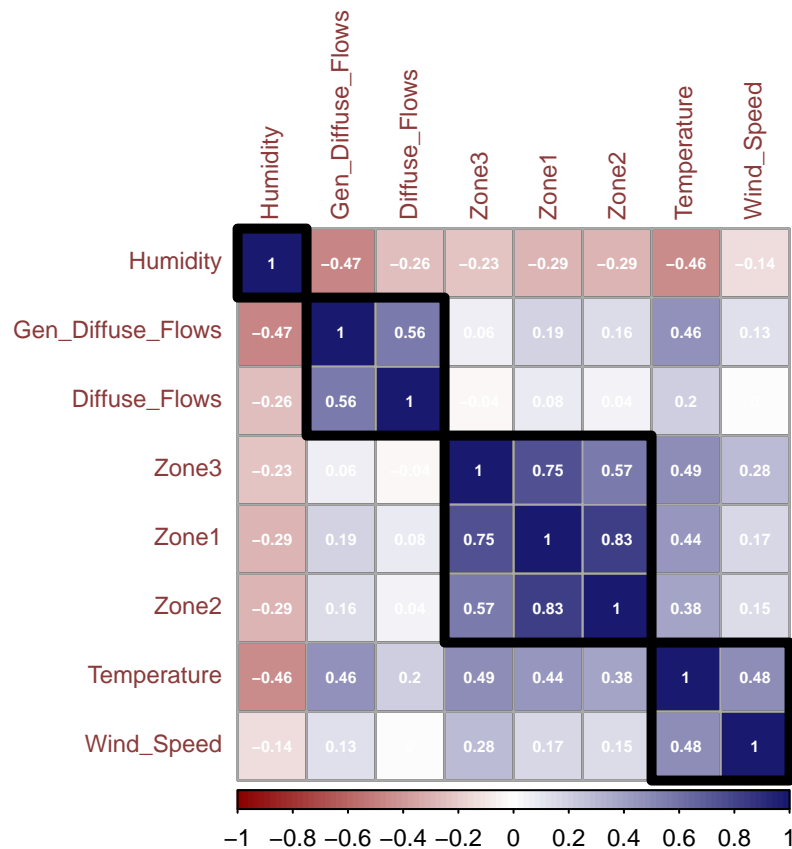
```
#Histograms of all variables
hist.data.frame(power[,2:9])
```



## Correlations

All three zones have strong correlations between each other. Temperature and wind speed have a correlation of 0.48. Humidity and temperature have a moderate but noticeable correlation as well. General diffuse flows and Diffuse flows have a correlation of 0.56. Temperature seems to show highest correlation with power consumption.

```
corr <- round(x = cor(power[, 2:9]), digits = 2)
corrplot(corr, method = "color", outline = T, addgrid.col = "darkgray",
  order="hclust", addrect = 4, rect.col = "black", rect.lwd = 5,
  cl.pos = "b", tl.col = "indianred4", tl.cex = .75, cl.cex = .75,
  addCoef.col = "white", number.digits = 2, number.cex = 0.5, col =
  colorRampPalette(c("darkred", "white", "midnightblue"))(100))
```

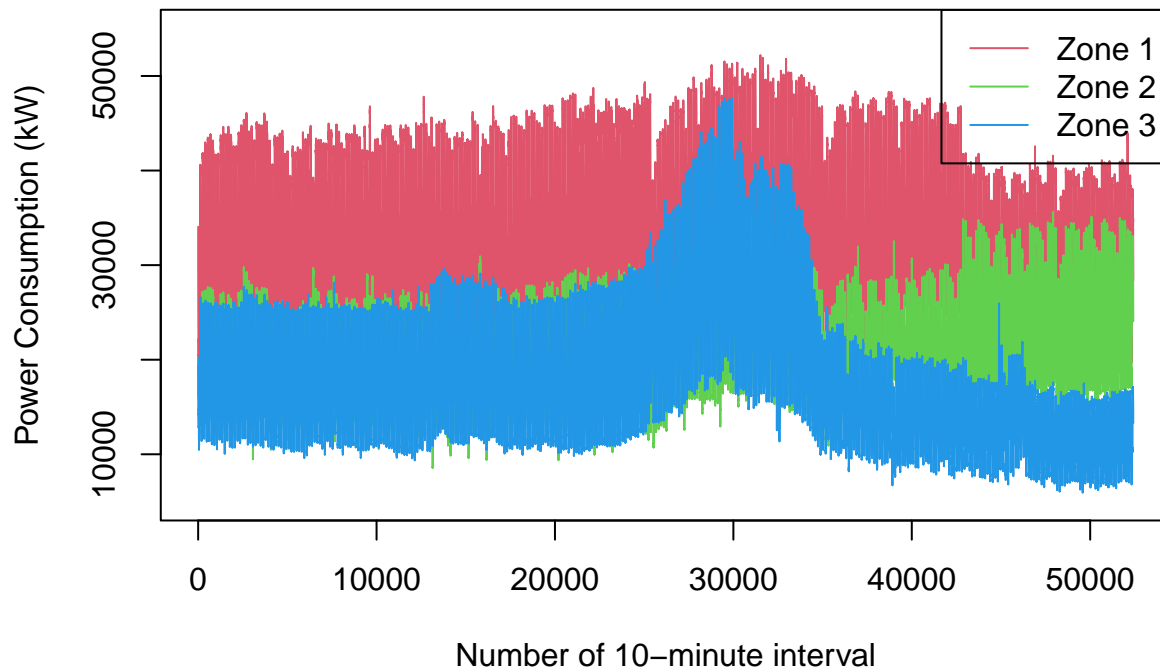


```
#Observe the the pair plots between variables
#plot(power[, 2:9], pch = 16, cex= 0.3)
```

### Comparisons of the Power Consumption Zones

Looking at all three zone's power consumption, it looks like Zone 1 has more power consumption than both Zone 2 and Zone 3.

```
plot(power$Zone1,
      type="l",
      col =2,
      ylim = c(5000,55000),
      xlab = "Number of 10-minute interval",
      ylab = "Power Consumption (kW)")
lines(power$Zone2,
      type="l",
      col =3)
lines(power$Zone3,
      type="l",
      col =4)
legend("topright", c("Zone 1", "Zone 2", "Zone 3"), lty = 1, col = 2:4)
```



## Explore Power Consumption in different time intervals

```
#Hourly Power consumption
power_hourly <- power %>%
  group_by(Hour= format(DateTime, "%Y-%m-%d %H")) %>%
  summarise(Total= sum(Zone1))

#Convert Zone1_Consumption to time series
hour.zone1.ts <- ts(power_hourly$Total, start= c(2017,1), end=c(2017, 8376), frequency=8376)
#Plot the ts
p1 <- autoplot(hour.zone1.ts, color="blue", main = "Hourly Power Consumption", cex =0.3)

#Daily Power consumption
power$Date <- as.Date(power$DateTime, format = "%m/%d/%Y")
power.daily.z1 <- power %>% group_by(Date) %>% summarise(Total= sum(Zone1))

p2 <- ggplot(power.daily.z1, aes(x = Date, y = Total)) +
  geom_line(size = 0.5, color = "blue") +
  labs(title = "Daily Power Consumption in Zone 1",
       x = "Day",
       y = "Power Consumption") +
  theme_classic()+
  theme(plot.title=element_text(hjust=0.3, size=10))

#Weekly Power consumption
power_weekly <- power %>%
  group_by(week = lubridate::week(Date)) %>% summarise(Total= sum(Zone1))

p3 <- ggplot(power_weekly, aes(x = week, y = Total)) +
  geom_line(size = 0.5, color = "blue") +
  labs(title = "Weekly Power Consumption in Zone 1",
```


```

    x = "Week",
    y = "Power Consumption") +
  theme_classic()+
  theme(plot.title=element_text(hjust=0.3, size=10))

#Monthly Power consumption
power_monthly <- power %>%
  group_by(Month = format(Date, "%Y-%m")) %>%
  summarise(Total = sum(Zone1))

p4 <- ggplot(power_monthly, aes(x = Month, y = Total, group = 1)) +
  geom_line(size = 0.5, color = "blue") +
  labs(title = "Monthly Power Consumption in Zone 1",
    x = "Month",
    y = "Power Consumption") +
  theme_classic()+
  theme(plot.title=element_text(hjust=0.3, size=10))+
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
grid.arrange(p1, p2, p3, p4, ncol = 2)

```




It is more interesting to explore the daily data, since we would be more interesting to forecast the power each Zone needed daily, so the power providers could increase or decrease power supply to meet the demand.

### Daily Power Consumption for all three zones

All three zones for power consumption have similar trending paths. The power consumption increases up until it reaches a peak followed by a decline. There appears to be some seasonality. There are some sharp

declines on random days, perhaps indicating a presence of power outages. Zone 1 and Zone 2 seem to have big daily fluctuation while Zone 3 seem less fluctuated. Also, They all seem to peak at around May to September which could be due to more power consumption during Summer time.

```
power$Date <- as.Date(power$DateTime, format = "%m/%d/%Y")
power.daily.z1 <- power %>% group_by(Date) %>% summarise(Total= sum(Zone1))
power.daily.z2 <- power %>% group_by(Date) %>% summarise(Total= sum(Zone2))
power.daily.z3 <- power %>% group_by(Date) %>% summarise(Total= sum(Zone3))
p1<- ggplot(power.daily.z1, aes(x = Date, y = Total)) +
  geom_line(size = 0.5, color = "blue") +
  labs(title = "Daily Power Consumption in Zone 1",
       x = "Daily",
       y = "Power Consumption") +
  theme_classic()+
  theme(plot.title=element_text(hjust=0.5))
p2 <- ggplot(power.daily.z2, aes(x = Date, y = Total)) +
  geom_line(size = 0.5, color = "blue") +
  labs(title = "Daily Power Consumption in Zone 2",
       x = "Daily",
       y = "Power Consumption") +
  theme_classic()+
  theme(plot.title=element_text(hjust=0.5))
p3 <- ggplot(power.daily.z3, aes(x = Date, y = Total)) +
  geom_line(size = 0.5, color = "blue") +
  labs(title = "Daily Power Consumption in Zone 3",
       x = "Daily",
       y = "Power Consumption") +
  theme_classic()+
  theme(plot.title=element_text(hjust=0.5))
grid.arrange(p1, p2, p3, ncol = 1)
```



#### Check for random walk


```
#Convert daily power in 3 zones into time series object
power.daily.z1.ts <- ts(power.daily.z1$Total)
power.daily.z2.ts <- ts(power.daily.z2$Total)
power.daily.z3.ts <- ts(power.daily.z3$Total)
```

Autocorrelation show high lag 1 correlation in all three zones with zone 1 very close to random walk while zone 2 and 3 also show correlation at multiple lags which could be weekly seasonality.


```
#Compute and plot the autocorrelation and produce an ACF plot
par(mfrow=c(2,2))
par(mar=c(5,5,5,2))
#To adjust the plot size and not get the title cut off
acf_value1 <- acf(power.daily.z1.ts, lag.max=364,
                  main='Zone 1 Autocorrelation Plot with Lag-1')
acf_value2 <- acf(power.daily.z2.ts, lag.max=364,
                  main='Zone 2 Autocorrelation Plot with Lag-1')
acf_value3 <- acf(power.daily.z3.ts, lag.max=364,
                  main='Zone 3 Autocorrelation Plot with Lag-1')
```




### Zone 1 Autocorrelation Plot with Lag



### Zone 2 Autocorrelation Plot with Lag



### Zone 3 Autocorrelation Plot with Lag



```
#Get value at lag-1  
print(acf_value1[1])
```

```
##  
## Autocorrelations of series 'power.daily.z1.ts', by lag  
##  
##      1  
## 0.875
```

```
print(acf_value2[1])
```

```
##  
## Autocorrelations of series 'power.daily.z2.ts', by lag  
##  
##      1  
## 0.866
```

```
print(acf_value3[1])
```

```
##  
## Autocorrelations of series 'power.daily.z3.ts', by lag  
##  
##      1  
## 0.989
```

### Examine the first Three Weeks of the Different Zones

The following shows the first few weeks of power consumption for all the zones. Zone 1 and 2 starts low on Sunday, goes high during the weekdays, and then low again on Saturday. Zone 3 seems to have the opposite behavior and the difference between days are less pronounced.

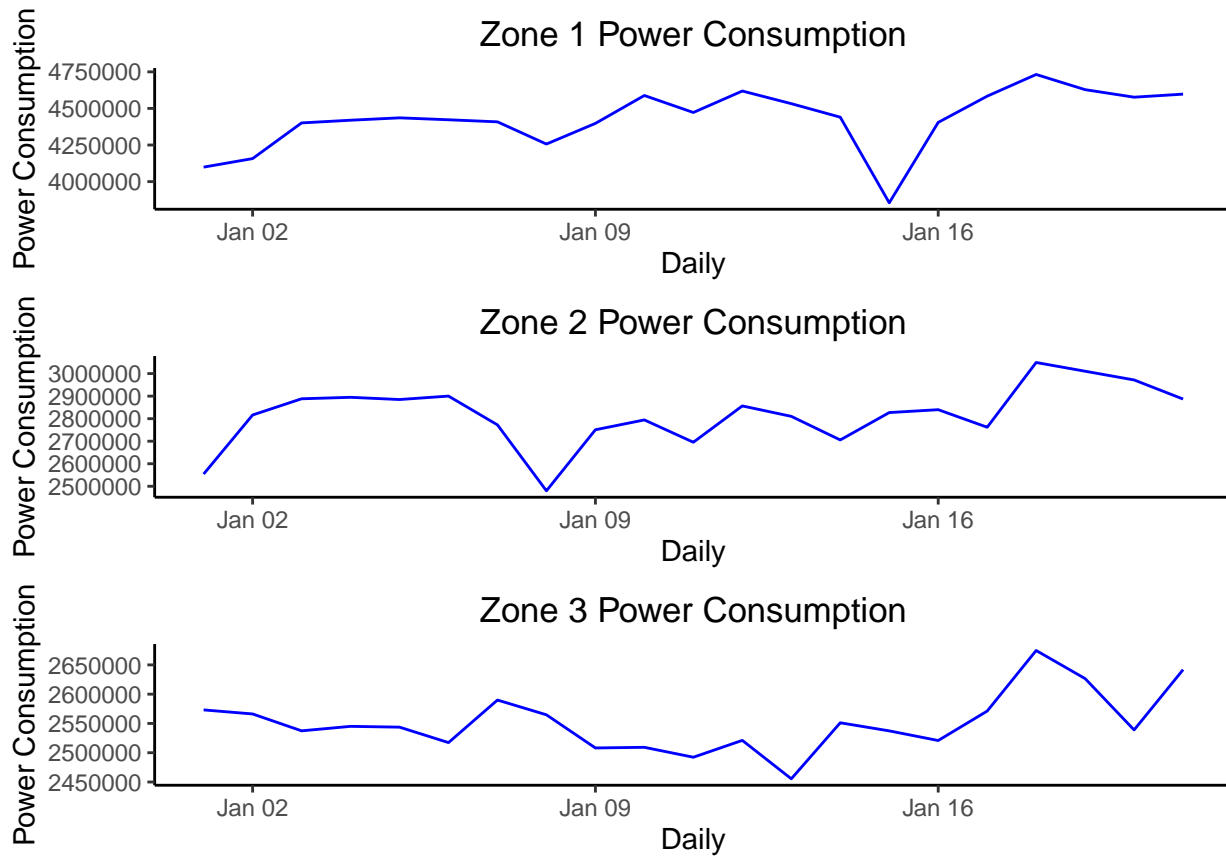
```

#Zoom in the first 3 weeks of Zone 1
p1 <- ggplot(power.daily.z1[0:21, ], aes(x = Date, y = Total)) +
  geom_line(size = 0.5, color = "blue") +
  labs(title = "Zone 1 Power Consumption",
       x = "Daily",
       y = "Power Consumption") +
  theme_classic()+
  theme(plot.title=element_text(hjust=0.5))

#Zoom in the first 3 weeks of Zone 1
p2 <- ggplot(power.daily.z2[0:21, ], aes(x = Date, y = Total)) +
  geom_line(size = 0.5, color = "blue") +
  labs(title = "Zone 2 Power Consumption",
       x = "Daily",
       y = "Power Consumption") +
  theme_classic()+
  theme(plot.title=element_text(hjust=0.5))

#Zoom in the first 3 weeks of Zone 3
p3 <- ggplot(power.daily.z3[0:21, ], aes(x = Date, y = Total)) +
  geom_line(size = 0.5, color = "blue") +
  labs(title = "Zone 3 Power Consumption",
       x = "Daily",
       y = "Power Consumption") +
  theme_classic()+
  theme(plot.title=element_text(hjust=0.5))
grid.arrange(p1, p2, p3, ncol = 1)

```



## Partitioning the Data

```
#Convert power data frame to daily data including all the variables
daily.sum <- power[,2:10] %>%
  group_by(Date) %>%
  summarise(across(everything(), sum), .groups = 'drop')
#Compute daily mean of the variables (Temperature, Humidity, Wind_Speed, Gen_Diffuse_Flows, Diffuse_Flows)
daily.var.mean <- (daily.sum[,2:6])/144
#Create day of the week column
daily.sum$DOW <- wday(daily.sum$Date, label = TRUE)
#create new frames with values DOW converted to dummies
DOW.dummies <- model.matrix(~ 0 + DOW, data = daily.sum)
#Rename each dummies column without "Dow" in front
colnames(DOW.dummies) <- gsub("DOW", "", colnames(DOW.dummies))
#Combine the data frames, exclude Tues to avoid dummy trap
X <- as.data.frame(cbind(daily.var.mean, DOW.dummies[, -3]))

#Create y series with power values from daily data frame for zone 1
y1 <- daily.sum$Zone1
y2 <- daily.sum$Zone2
y3 <- daily.sum$Zone3
#Split data
nTotal <- length(y1) #could use length zone 1 for all 3 zones
nValid <- 56
nTrain <- nTotal - nValid
xTrain <- X[1:nTrain, ]
```

```

xValid <- X[(nTrain + 1):nTotal, ]
#
##Response variable for train and valid sets
# Zone 1
yTrain1 <- y1[1:nTrain]
yValid1 <- y1[(nTrain + 1):nTotal]

yTrain2 <- y2[1:nTrain]
yValid2 <- y2[(nTrain + 1):nTotal]

yTrain3 <- y3[1:nTrain]
yValid3 <- y3[(nTrain + 1):nTotal]

#Convert y train and validation into time series object
yTrain1.ts <- ts(yTrain1, start = c(1, 1), end= c(44, 7), frequency = 7)
yValid1.ts <- ts(yValid1, start = c(45, 1), end= c(52, 7), frequency = 7)

yTrain2.ts <- ts(yTrain2, start = c(1, 1), end= c(44, 7), frequency = 7)
yValid2.ts <- ts(yValid2, start = c(45, 1), end= c(52, 7), frequency = 7)

yTrain3.ts <- ts(yTrain3, start = c(1, 1), end= c(44, 7), frequency = 7)
yValid3.ts <- ts(yValid3, start = c(45, 1), end= c(52, 7), frequency = 7)

```


## Decomposition of the Time Series

```


#Decompose time series
stl1.run <- stl(yTrain1.ts, s.window = "periodic")
stl2.run <- stl(yTrain2.ts, s.window = "periodic")
stl3.run <- stl(yTrain3.ts, s.window = "periodic")

plot(stl1.run)


```



```
plot(stl2.run)
```



```
plot(stl3.run)
```



```
# autoplot(yTrain1.ts ) + autolayer(yValid1.ts, color = "blue") + theme_classic() +
#   labs(title = "Training and Validation Sets", y = "Power Consumption (kW)")
# autoplot(yTrain2.ts ) + autolayer(yValid2.ts, color = "blue") + theme_classic() +
#   labs(title = "Training and Validation Sets", y = "Power Consumption (kW)")
# autoplot(yTrain3.ts ) + autolayer(yValid3.ts, color = "blue") + theme_classic() +
#   labs(title = "Training and Validation Sets", y = "Power Consumption (kW)")
```

## Modeling

### Mean Model

```
#Average model function
mean_predict <- function(yTrain.ts, yValid.ts, titl) {
  mean_power = meanf(yTrain.ts, h=nValid )
  mean_power.ts <- ts( mean_power$mean, start = c(45, 1), end= c(52, 7),
    frequency = 7)
  p <- autoplot(yTrain.ts) +
    autolayer(mean_power$mean, color = 'red')+
    autolayer(yValid.ts, color = "blue")+
    labs(title =titl,x = "week", y = "Power Consumption")+
    theme (title =element_text(size=7))+
    geom_hline(yintercept = mean_power$mean, color = "green", size = 0.5)
  acc <- accuracy(mean_power$mean, yValid.ts)
  lst <- list(acc,p)
  return(lst)}
#call model output
zone1.mean<- mean_predict(yTrain1.ts, yValid1.ts, "Zone 1")
zone2.mean<- mean_predict(yTrain2.ts, yValid2.ts, "Zone 2")
zone3.mean<- mean_predict(yTrain3.ts, yValid3.ts, "Zone 3")
```

```
zone1.mean[[1]]
```

```
##           ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set -573375.5 586831.7 573375.5 -13.84628 13.84628 0.2941203  3.904904
```

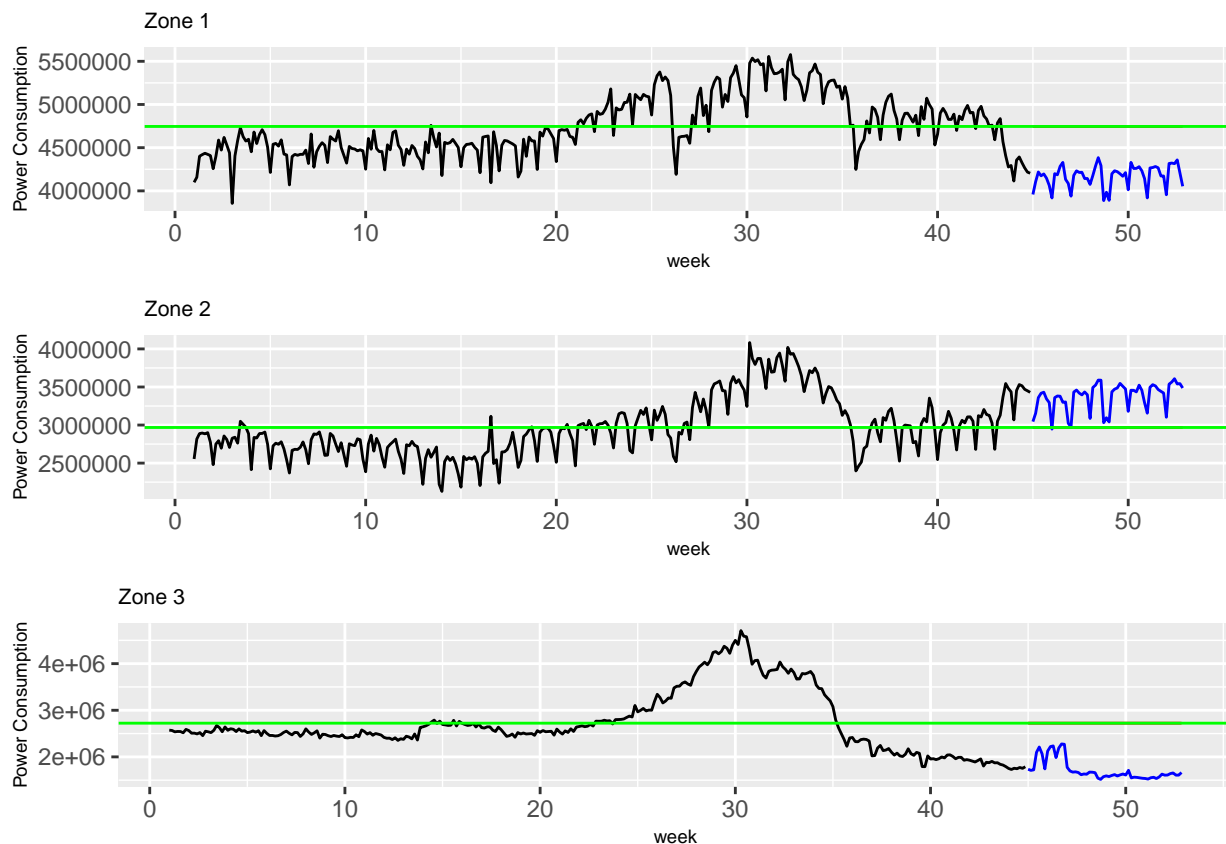
```
zone2.mean[[1]]
```

```
##           ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set 404226.2 441012.8 404983.5 11.73235 11.75804 0.3748593  2.172706
```

```
zone3.mean[[1]]
```

```
##           ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set -1004672 1028633 1004672 -60.70393 60.70393 0.8385883  9.670353
```

```
ggarrange(zone1.mean[[2]], zone2.mean[[2]], zone3.mean[[2]], ncol=1)
```



## Naive Forecast

```
# Naive Forecast function for all 3 zones
naive <- function(yTrain.ts, yValid.ts, titl) {
  naive_model = forecast::naive(yTrain.ts, h=nValid)
  #naive_model.pred<- forecast(naive_model, h = nValid, level = c(.95))
  naive_model.ts <- ts(naive_model$mean, start = c(45, 1), end= c(52, 7),
    frequency = 7)
  p <- autoplot(yTrain.ts) +
    autolayer(naive_model$mean, color = 'red')+
    autolayer(yValid.ts, color = "blue")+
    labs(title =titl,x = "week", y = "Power Consumption")+

```

```

  theme (title =element_text(size=7))
  #geom_hline(yintercept = (naive_model$mean, color = "green", size = 0.5))
  acc <- accuracy(naive_model$mean, yValid.ts)
  lst <- list(acc,p)
  return(lst)}
#call model output
zone1.naive<- naive(yTrain1.ts, yValid1.ts, "Zone 1")
zone2.naive<- naive(yTrain2.ts, yValid2.ts, "Zone 2")
zone3.naive<- naive(yTrain3.ts, yValid3.ts, "Zone 3")

```

```
zone1.naive[[1]]
```

```
##              ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set -31758.88 128921.2 95225.98 -0.8537643 2.32781 0.2941203 0.8467079
```


```
zone2.naive[[1]]
```

```
##              ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set -57086.56 185343.7 130814.6 -1.987304 4.073163 0.3748593 0.9106244
```

```
zone3.naive[[1]]
```

```
##              ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set -70039.93 231569.8 205595.1 -5.54176 11.74121 0.8385883 2.004269
```

```
ggarrange(zone1.naive[[2]], zone2.naive[[2]], zone3.naive[[2]], ncol=1)
```





## Seasonal Naive Forecast

```
# Build seasonal Naive Forecast function for all 3 zones
snaive <- function(yTrain.ts, yValid.ts, titl) {
  snaive_model = forecast::snaive(yTrain.ts, h=nValid)
  snaive_model.ts <- ts(snaive_model$mean, start = c(45, 1), end= c(52, 7),
                        frequency = 7)

  p <- autoplot(yTrain.ts) +
    autolayer(snaive_model$mean, color = 'red')+
    autolayer(yValid.ts, color = "blue")+
    labs(title =titl,x = "week", y = "Power Consumption")+
    theme (title =element_text(size=7))
    #geom_hline(yintercept = (naive_model$mean, color = "green", size = 0.5))
  acc <- accuracy(snaive_model$mean, yValid.ts)
  lst <- list(acc,p)
  return(lst)}

#call model output
zone1.snaive<- snaive(yTrain1.ts, yValid1.ts, "Zone 1")
zone2.snaive<- snaive(yTrain2.ts, yValid2.ts, "Zone 2")
zone3.snaive<- snaive(yTrain3.ts, yValid3.ts, "Zone 3")

zone1.snaive[[1]]

##           ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set -95450.4 127311.1 105459.7 -2.332443 2.563294 0.3962262 0.857502


zone2.snaive[[1]]

##           ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set -44567.12 125935.4 87784.12 -1.432863 2.683949 0.4514328 0.6385044

zone3.snaive[[1]]

##           ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set -38740.71 224690 186269.9 -3.69554 10.4615 0.8313098 1.890542

ggarrange(zone1.snaive[[2]], zone2.snaive[[2]], zone3.snaive[[2]], ncol=1)
```



## Holt Winters

```
#Built function for Holt Winters model
holtW<- function(yTrain.ts, yValid.ts, titl) {
  holt.model <- ets(yTrain.ts, model = "ZAA", alpha = .2, gamma = .05)
  holt.model.pred <- forecast(holt.model, h = nValid)
  holt.model.pred.ts <- ts(holt.model.pred$mean, start = c(45, 1), end= c(52, 7), frequency = 7)
  model <- holt.model
  p <- autoplot(yTrain.ts) +
    autolayer( holt.model.pred, color = 'red')+
    autolayer(yValid.ts, color = "blue")+
    labs(title =titl,x = "week", y = "Power Consumption")+
    theme (title =element_text(size=7))
  acc <- accuracy(holt.model.pred, yValid.ts)
  lst <- list(acc,p, model)
  return(lst)}

#call model output
zone1.holt<- holtW(yTrain1.ts, yValid1.ts, "Zone 1")
zone2.holt<- holtW(yTrain2.ts, yValid2.ts, "Zone 2")
zone3.holt<- holtW(yTrain3.ts, yValid3.ts, "Zone 3")

zone1.holt[[1]]
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -1564.421 151970.29 101378.05 -0.1093762 2.156048 0.6155733
## Test set      8045.769  76833.47  57968.65  0.1542896 1.391692 0.3519890
##              ACF1 Theil's U
```

```
## Training set 0.4544233      NA
## Test set      0.4423987 0.5159244
```


```
zone2.holt[[1]]
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set   5126.657 133642.4  90436.87  0.01413649  3.077895  0.5900849
## Test set      -101729.396 151507.5 112031.94 -3.14159657  3.434808  0.7309890
##              ACF1 Theil's U
## Training set   0.4763998      NA
## Test set       0.5174548 0.7628992
```

```
zone3.holt[[1]]
```


```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  -2443.093  96697.0  63366.07 -0.1357273  2.290272  0.4209035
## Test set      -78124.597 238432.9 213787.46 -6.0393713 12.245072  1.4200642
##              ACF1 Theil's U
## Training set   0.6069677      NA
## Test set       0.8287440  2.077576
```

```
ggarrange(zone1.holt[[2]], zone2.holt[[2]], zone3.holt[[2]], ncol=1)
```



```
#Residuals plot for Holt model
plot(zone1.holt[[3]]$residuals, xlab = "Year", ylab = "Residuals", bty = "l",
      lwd = 1, main = "Residuals of Zone 1 power consumption- Holt's Model")
```

## Residuals of Zone 1 power consumption– Holt's Model



## TSLM

```
#Built function for TSLM
tslm<- function(yTrain.ts, yValid.ts, titl) {
  #Create the formula for the regression model
  (formula <- as.formula(paste("yTrain1.ts", paste(c("trend", "season",
                                                    colnames(xTrain)),
                                                    collapse = "+"), sep = "~")))

  tslm.model <- forecast::tslm(formula, data = xTrain)
  tslm.model.pred <- forecast(tslm.model, newdata = xValid)
  tslm.model.pred.ts <- ts(tslm.model.pred$mean, start = c(45, 1), end= c(52, 7), frequency = 7)
  p <- autoplot(yTrain.ts) +
    autolayer( tslm.model.pred, color = 'red')+
    autolayer(yValid.ts, color = "blue")+
    labs(title =titl,x = "week", y = "Power Consumption")+
    theme (title =element_text(size=7))
  acc <- accuracy(tslm.model.pred, yValid.ts)
  lst <- list(acc,p)
  return(lst)}

#call model output
zone1.tslm<- tslm(yTrain1.ts, yValid1.ts, "Zone 1")

## Warning in predict.lm(predict_object, newdata = newdata, se.fit = TRUE, :
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(predict_object, newdata = newdata, se.fit = TRUE, :
## prediction from a rank-deficient fit may be misleading
```

```

zone2.tslm<- tslm(yTrain2.ts, yValid2.ts, "Zone 2")

## Warning in predict.lm(predict_object, newdata = newdata, se.fit = TRUE, :
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(predict_object, newdata = newdata, se.fit = TRUE, :
## prediction from a rank-deficient fit may be misleading

zone3.tslm<- tslm(yTrain3.ts, yValid3.ts, "Zone 3")

## Warning in predict.lm(predict_object, newdata = newdata, se.fit = TRUE, :
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(predict_object, newdata = newdata, se.fit = TRUE, :
## prediction from a rank-deficient fit may be misleading

zone1.tslm[[1]]

##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  2.421659e-11 214543.3 168725.1 -0.2015387 3.569670 1.024508
## Test set     -2.389210e+05 298700.5 247018.1 -5.7833709 5.973748 1.499908
##              ACF1 Theil's U
## Training set  0.7255617      NA
## Test set      0.6972350  2.001817

zone2.tslm[[1]]

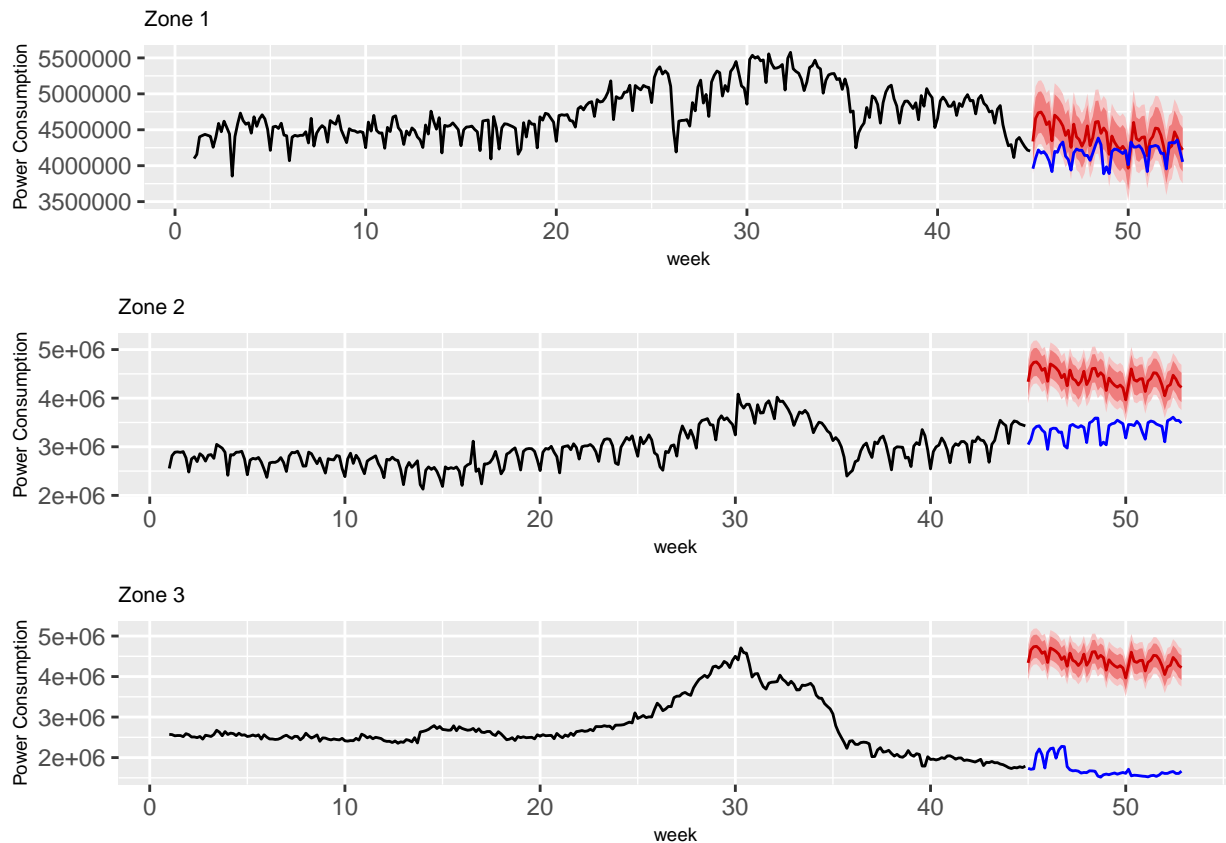
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  2.421659e-11 214543.3 168725.1 -0.2015387 3.56967 1.024508
## Test set     -1.039274e+06 1064187.8 1039274.0 -31.1550282 31.15503 6.310531
##              ACF1 Theil's U
## Training set  0.7255617      NA
## Test set      0.7220227  5.294553

zone3.tslm[[1]]

##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  2.421659e-11 214543.3 168725.1 -0.2015387 3.56967 1.024508
## Test set     -2.693242e+06 2701385.0 2693241.5 -159.8616239 159.86162 16.353516
##              ACF1 Theil's U
## Training set  0.7255617      NA
## Test set      0.4921255  24.84325

ggarrange(zone1.tslm[[2]], zone2.tslm[[2]], zone3.tslm[[2]], ncol=1)

```



## Auto Arima Model

```
#Built function for TSLM model in 3 zones
autoArima<- function(yTrain.ts, yValid.ts, titl) {
  predictor_train <- as.matrix(xTrain)
  predictors_test <- as.matrix(xValid)
  autoArima.model <- forecast::auto.arima(yTrain1.ts, xreg =predictor_train)
  autoArima.model.pred <- forecast(autoArima.model, h = nValid, xreg= predictors_test)
  autoArima.model.pred.ts <- ts(autoArima.model.pred$mean, start = c(45, 1), end= c(52,7), frequency = 7)
  model <- autoArima.model
  p <- autoplot(yTrain.ts) +
    autolayer( autoArima.model.pred, color = 'red')+
    autolayer(yValid.ts, color = "blue")+
    labs(title =titl,x = "week", y = "Power Consumption")+
    theme (title =element_text(size=7))
  acc <- accuracy(autoArima.model.pred, yValid.ts)
  lst <- list(acc,p,model)
  return(lst)}

#call model output
zone1.autoArima<- autoArima(yTrain1.ts, yValid1.ts, "Zone 1")
zone2.autoArima<- autoArima(yTrain2.ts, yValid2.ts, "Zone 2")
zone3.autoArima<- autoArima(yTrain3.ts, yValid3.ts, "Zone 3")

#Output auto arima performance
zone1.autoArima[[1]]
```

##	ME	RMSE	MAE	MPE	MAPE	MASE
----	----	------	-----	-----	------	------

```
## Training set      3100.921 130277.8 88525.01 -0.01165145 1.887858 0.5375289
## Test set         -260420.338 274318.3 260420.34 -6.26871282 6.268713 1.5812871
##                   ACF1 Theil's U
## Training set     -0.008193179      NA
## Test set         0.456826861    1.84643
```


```
zone2.autoArima[[1]]
```

```
##                   ME      RMSE      MAE      MPE      MAPE      MASE
## Training set      3100.921 130277.8 88525.01 -0.01165145 1.887858 0.5375289
## Test set         -1060773.339 1065558.5 1060773.34 -31.66728795 31.667288 6.4410762
##                   ACF1 Theil's U
## Training set     -0.008193179      NA
## Test set         0.349214670  5.291881
```

```
zone3.autoArima[[1]]
```

```
##                   ME      RMSE      MAE      MPE      MAPE
## Training set      3100.921 130277.8 88525.01 -0.01165145 1.887858
## Test set         -2714740.835 2730131.0 2714740.84 -161.93420553 161.934206
##                   MASE      ACF1 Theil's U
## Training set      0.5375289 -0.008193179      NA
## Test set         16.4840611  0.798755988 25.37145
```

```
ggarrange(zone1.autoArima[[2]], zone2.autoArima[[2]], zone3.autoArima[[2]], ncol=1)
```



```
#Output auto arima parameters
```

```
zone1.autoArima[[3]]
```

```
## Series: yTrain1.ts
```

```
## Regression with ARIMA(2,0,1)(2,0,1)[7] errors
##
## Coefficients:
##          ar1      ar2      ma1      sar1      sar2      sma1  intercept
##          1.3438 -0.3589 -0.6620 -0.6534 -0.0508  0.7051 4561988.8
## s.e.      0.1925  0.1832  0.1638  0.2624  0.0713  0.2591 225397.9
##          Temperature Humidity Wind_Speed Gen_Diffuse_Flows Diffuse_Flows
##          1796.499  1096.416  1847.980          190.3560        -89.6163
## s.e.      6503.588  1005.791  6224.514          169.1815        298.2813
##          Sun      Mon      Wed      Thu      Fri      Sat
##          -277573.6 -10841.20 25232.93 19962.89 -42749.64 -52445.9
## s.e.      24722.6  20558.59 20566.58 24826.10 26363.81 26395.2
##
## sigma^2 = 1.803e+10: log likelihood = -4065.63
## AIC=8169.26 AICc=8171.9 BIC=8240.13
```

```
zone2.autoArima[[3]]
```

```
## Series: yTrain1.ts
## Regression with ARIMA(2,0,1)(2,0,1)[7] errors
##
## Coefficients:
##          ar1      ar2      ma1      sar1      sar2      sma1  intercept
##          1.3438 -0.3589 -0.6620 -0.6534 -0.0508  0.7051 4561988.8
## s.e.      0.1925  0.1832  0.1638  0.2624  0.0713  0.2591 225397.9
##          Temperature Humidity Wind_Speed Gen_Diffuse_Flows Diffuse_Flows
##          1796.499  1096.416  1847.980          190.3560        -89.6163
## s.e.      6503.588  1005.791  6224.514          169.1815        298.2813
##          Sun      Mon      Wed      Thu      Fri      Sat
##          -277573.6 -10841.20 25232.93 19962.89 -42749.64 -52445.9
## s.e.      24722.6  20558.59 20566.58 24826.10 26363.81 26395.2
##
## sigma^2 = 1.803e+10: log likelihood = -4065.63
## AIC=8169.26 AICc=8171.9 BIC=8240.13
```


```
zone3.autoArima[[3]]
```

```
## Series: yTrain1.ts
## Regression with ARIMA(2,0,1)(2,0,1)[7] errors
##
## Coefficients:
##          ar1      ar2      ma1      sar1      sar2      sma1  intercept
##          1.3438 -0.3589 -0.6620 -0.6534 -0.0508  0.7051 4561988.8
## s.e.      0.1925  0.1832  0.1638  0.2624  0.0713  0.2591 225397.9
##          Temperature Humidity Wind_Speed Gen_Diffuse_Flows Diffuse_Flows
##          1796.499  1096.416  1847.980          190.3560        -89.6163
## s.e.      6503.588  1005.791  6224.514          169.1815        298.2813
##          Sun      Mon      Wed      Thu      Fri      Sat
##          -277573.6 -10841.20 25232.93 19962.89 -42749.64 -52445.9
## s.e.      24722.6  20558.59 20566.58 24826.10 26363.81 26395.2
##
## sigma^2 = 1.803e+10: log likelihood = -4065.63
## AIC=8169.26 AICc=8171.9 BIC=8240.13
```




## Optimized ARIMA model on Zone 1

```
#Plot autocorrelation of yTrain1.ts  
ggtsdisplay(yTrain1.ts)
```



```
#Plot autocorrelation of differenced yTrain1.ts after two non-seasonal differences and 1 seasonal difference  
yTrain1.ts %>% diff %>% diff %>% diff(lag = 7) %>% ggtsdisplay(lag=44)
```



After taking two non-seasonal differences and 1 seasonal difference, the autocorrelation seems improved and has removed some of the seasonality.

```
#Built function for my Arima models in 3 zones
myArima<- function(yTrain.ts, yValid.ts, titl) {
  predictor_train <- as.matrix(xTrain)
  predictors_test <- as.matrix(xValid)
  myArima.model <- forecast::Arima(yTrain1.ts, order = c(1,2,1),
                                   seasonal=c(1,0,0), xreg =predictor_train)
  myArima.model.pred <- forecast(myArima.model, h = nValid, xreg= predictors_test)
  myArima.model.pred.ts <- ts(myArima.model.pred$mean, start = c(45, 1), end= c(52,7), frequency = 7)
  model <- myArima.model
  p <- autoplot(yTrain.ts, series='Train', color = "black") +
    autolayer( myArima.model.pred, color = 'red', series='Forecast')+
    autolayer(yValid.ts, color = "blue", series = 'Test')+
    labs(title =titl,x = "week", y = "Power Consumption")+
    theme (title =element_text(size=7))+
    guides(color = guide_legend(title = "Data Series")) +
    scale_color_manual(values = c(Train = "black", Forecast = "red",
                                   Test = "blue"))

  #coord_cartesian(xlim = c(15, 54))
  acc <- accuracy(myArima.model.pred, yValid.ts)
  lst <- list(acc,p, model)
  return(lst)}

#call model output
zone1.myArima<- myArima(yTrain1.ts, yValid1.ts, "Zone 1")
zone2.myArima<- myArima(yTrain2.ts, yValid2.ts, "Zone 2")
zone3.myArima<- myArima(yTrain3.ts, yValid3.ts, "Zone 3")
```

```
#Output auto arima performance
```

```
zone1.myArima[[1]]
```

```
##                ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -4872.857 133783.69 88708.39 -0.1518214 1.894779 0.5386424
## Test set     -32504.793 80930.79 61207.14 -0.8143030 1.484744 0.3716533
##                ACF1 Theil's U
## Training set -0.02427314      NA
## Test set      0.38306457 0.5460399
```


```
zone2.myArima[[1]]
```

```
##                ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -4872.857 133783.7 88708.39 -0.1518214 1.894779 0.5386424
## Test set     -832857.795 841332.3 832857.79 -24.9363272 24.936327 5.0571600
##                ACF1 Theil's U
## Training set -0.02427314      NA
## Test set      0.52757630 4.180559
```

```
zone3.myArima[[1]]
```

```
##                ME      RMSE      MAE      MPE      MAPE
## Training set -4872.857 133783.7 88708.39 -0.1518214 1.894779
## Test set     -2486825.290 2498228.4 2486825.29 -148.1704264 148.170426
##                MASE      ACF1 Theil's U
## Training set 0.5386424 -0.02427314      NA
## Test set     15.1001449 0.73316989 23.14676
```


```
grid.arrange(zone1.myArima[[2]], zone2.myArima[[2]], zone3.myArima[[2]], ncol=1)
```




## Residual plots for Holt's and Arima model

```
par(mfrow=c(2,1))
#Residuals plot for Holt model
plot(zone1.holt[[3]]$residuals, xlab = "Year", ylab = "Residuals", bty = "l",
     lwd = 1, main = "Residuals of Zone 1 power consumption- Holt's Model")
#Residuals plot for arima model
plot(zone1.myArima[[3]]$residuals, xlab = "Year", ylab = "Residuals", bty = "l",
     lwd = 1, main = "Residuals of Zone 1 power consumption- Arima Model")
```

### Residuals of Zone 1 power consumption– Holt's Model



### Residuals of Zone 1 power consumption– Arima Model



## Neuro Network

```
#Built function for my Arima models in 3 zones
nnetar<- function(yTrain.ts, yValid.ts, titl) {
  nnetar.model <- forecast::nnetar(yTrain1.ts)
  nnetar.model.pred <- forecast(nnetar.model, h = nValid)
  #nnetar.model.pred.ts <- ts(nnetar.model.pred, start = c(45, 1), end= c(52,7), frequency = 7)
  p <- autoplot(yTrain.ts, series='Train', color = "black") +
    autolayer( nnetar.model.pred, color = 'red', series='Forecast')+
    autolayer(yValid.ts, color = "blue", series = 'Test')+
    labs(title =titl,x = "week", y = "Power Consumption")+
    theme (title =element_text(size=7))+
    guides(color = guide_legend(title = "Data Series")) +
    scale_color_manual(values = c(Train = "black", Forecast = "red",
                                   Test = "blue"))

  #coord_cartesian(xlim = c(15, 54))
  acc <- accuracy(nnetar.model.pred, yValid.ts)
  lst <- list(acc,p)
  return(lst)}
```

```

#call model output
zone1.nnetar<- nnetar(yTrain1.ts, yValid1.ts, "Zone 1")
zone2.nnetar<- nnetar(yTrain2.ts, yValid2.ts, "Zone 2")
zone3.nnetar<- nnetar(yTrain3.ts, yValid3.ts, "Zone 3")

#Output auto arima performance
zone1.nnetar[[1]]

##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set    225.8251  54156.38  35419.41 -0.02358917  0.7549447  0.2150687
## Test set       -289644.0365 313655.66 289644.04 -7.02972900  7.0297290  1.7587351
##              ACF1 Theil's U
## Training set  0.07300538      NA
## Test set      0.31786933  2.118797

zone2.nnetar[[1]]


##              ME      RMSE      MAE      MPE      MAPE
## Training set   -108.2002   52347.5   33996.09 -0.0303586  0.7224266
## Test set       -1091461.1810 1104931.1 1091461.18 -32.7190310 32.7190310
##              MASE      ACF1 Theil's U
## Training set  0.2064262 0.03018241      NA
## Test set      6.6274145 0.41419058  5.505867

zone3.nnetar[[1]]

##              ME      RMSE      MAE      MPE      MAPE
## Training set    178.2742   50509.29   33859.56 -0.02358424  0.7190109
## Test set       -2751192.6832 2759897.98 2751192.68 -163.73761505 163.7376151
##              MASE      ACF1 Theil's U
## Training set  0.2055972 0.1017592      NA
## Test set      16.7053988 0.7822626  25.52062


grid.arrange(zone1.nnetar[[2]], zone2.myArima[[2]], zone3.nnetar[[2]], ncol=1)

```



Provide two weeks future forecast

```
fc.period = 14
#Convert full data to ts
y1.ts <- ts(y1, start = c(1, 1), end= c(52, 7), frequency = 7)
full.holt.model <- ets(y1.ts, model = "ZAA", alpha = .2, gamma = .05)
full.holt.model.pred <- forecast(full.holt.model, h = fc.period)
full.holt.model.pred.ts <- ts(full.holt.model.pred$mean, start = c(45, 1), end= c(52, 7), frequency = 7)
autoplot(y1.ts) +
  autolayer( full.holt.model.pred, color = 'orange')+
  labs(title = "Two week Forecast Plot",x = "week", y = "Power Consumption")+
  theme (title =element_text(size=7))
```



```
#Forecast Value  
full.holt.model.pred$mean
```

```
## Time Series:  
## Start = c(53, 1)  
## End = c(54, 7)  
## Frequency = 7  
## [1] 3998163 4254857 4274085 4275799 4269856 4176025 4165649 3987405 4246251  
## [10] 4267200 4270291 4265449 4172500 4162828
```